

Hack&Slash: Gesture Recognition for Dashboards

Seminar Thesis

Nico Rausch

Richard Tran

Giorgio Groß

1856484 - 1835121 - 1938927

At the Department of Economics and Management
Institute of Information Systems and Marketing (IISM)
Information & Market Engineering

Reviewer: Prof. Dr. Alexander Mädche
Advisor: Eric Pescera (TECO)

20th of July 2017

Contents

1. Introduction	1
2. Related Work	2
3. Dashboard	3
4. System Design	4
4.1. Components	4
4.1.1. Connector	5
4.1.2. Input	5
4.1.3. Recorder	5
4.1.4. Scanner	5
4.1.5. Dashboard	6
4.1.6. Recognition	6
5. Data Collection & Manipulation	7
5.1. Data Collection	7
5.2. Feature Design	7
5.2.1. Preprocessing	7
5.2.2. Features	7
5.3. Data Classification	8
Appendix	9
A. First Appendix Section	9

1. Introduction

Nowadays gesture recognition has achieved a high usage and meaning, especially for medical purposes and industrial production processes. These new technologies are available not only for industrial purposes, but also for the mass market, for which devices like the Xbox Kinect or the Wii Remote set the basis a few years ago. Until now, many research groups (**Zitation**) have already worked on technologies which enable the user to control computers with gestures. Most of these devices work vision-based, i.e. with cameras which recognize the visual changes caused by body movement and translate them into computer commands. These applications have the disadvantage that the user cannot move freely in the room, since he is bounded to the space captured by the camera in order to transmit the commands properly. This can be impractical for specific situations, e.g. presentations, when flexibility is needed considering the movement in the room.

ToDo

Furthermore, there are also several technologies which track the physical data of the movement with the help of sensors (accelerometer, gyroscope, etc.) to recognize a gesture. These can be used independently from the location of the user in the room, but often demands expensive equipment, for instance a data glove. In order to develop a cheap and accurate sensor based gesture recognition device, different approaches use data from an accelerometer to classify gestures. These show a high accuracy, but are limited in the number of different identifiable gestures due to the less data.

In our work, we used the “Thunderboard React”, which is a development platform with several sensors including an accelerometer and a gyroscope. We examined, how well such a sensor platform is suited for gesture recognition and if we could get even better results than previous works in terms of accuracy and number of recognizable gestures. We did this by calculating “features” like the maximal acceleration from the collected data and classified these with different machine learning approaches. Overall, we could recognize **4** different gestures with an average accuracy of **95**%. Furthermore, we created a dashboard to test the usability of our implementation.

ToDo

ToDo

2. Related Work

3. Dashboard

The main goal of our seminar work was to implement a gesture recognition program with which we can reliably control applications like dashboards on the computer. Therefore, we created a test dashboard with which we examined the usability of our program. It consists of several diagrams, which can be modified by steering elements like a dropdown menu and radio buttons. The main characteristic of the dashboard is that it can be controlled only with **ToDo** keyboard

Of course, this is just a simplified example of a dashboard with only a few features compared to what can be found in actual corporate dashboards, which allow more sophisticated commands. Nevertheless, our dashboard can be easily extended with further features by implementing more steering elements like slide bars or check buttons and by adding more recognizable gestures with which these elements can be controlled.




Figure 3.1.: Our demo application front end

4. System Design

In order to maximize reusability of code we decided to split the app into the six main components shown below. As you can see, our project consists of two individual parts, the recorder app and the recognizer app, sharing two helper modules, labeled Input and Connector. We implemented a core module for each part which is responsible to coordinate data streams and event responses. As such, the core modules instantiate their helper modules and use their interfaces to communicate with them. It is also worth knowing that some modules are divided into multiple files which are documented directly in our code.

4.1. Components




Figure 4.1.: Recognizer App




Figure 4.2.: Recorder App

4.1.1. Connector

As we ran into several issues trying to connect the Thunderboard React using Python we implemented the connector in JavaScript and used NodeJS libraries for interfacing the bluetooth device. To launch the recorder or recognizer app you need to launch the connector and pass it the path to the desired core module you aim to launch. After launching the connector it performs the steps shown in the following flow chart.

4.1.2. Input

The input module starts a new thread upon instantiation to avoid blocking the calling thread. In its new thread it keeps reading the sensor data from stdin, parses it into floating point values and inserts them into a ring buffer. Its interfaces can be used to set up the buffer, to reset the buffer, to insert new data into the buffer or to start recording. In the latter case, the module finds out when the actual gesture started on its own by comparing the acceleration vector length to a threshold value. As soon as the buffer fills up its content are written to a static CSV file autonomously. As you can see, the input module never takes action on its own but is still smart enough to perform a full control flow to record a gesture or to maintain the ring buffer.

4.1.3. Recorder

The recorder module provides a simple user interface which is used to specify the gesture you want to record as well as to explicitly start recording. That is done through a button which has to be pressed before each single gesture; we implemented it this way in order to avoid accidentally record a gesture which would pollute our training set. It also is responsible for instantiation and hooking up the input module.

4.1.4. Scanner

The scanner module is only used by the recognition module. It is used for monitoring the ring buffer content and checks whether a gesture was observed. Thus, it instantiates, trains

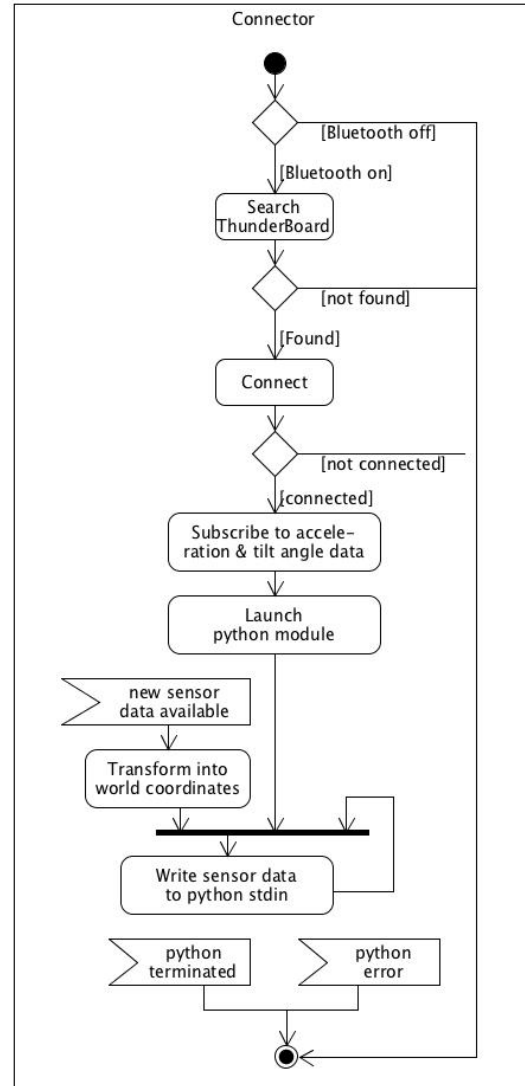


Figure 4.3.: Establishing the connection to the Thunderboard React

and maintains the machine learning algorithm used to classify live sensor data. Training includes transforming the recorded raw acceleration and tilt angle data into features and feeding them to the machine learning algorithm. All live sensor data will also have to be transformed into the same features so that it can be classified. We describe what features are and how we transformed the raw data in the next chapter. To avoid unnecessary complexity, we included the cross validation used to evaluate the performance of different machine learning algorithms directly into the scanner module. Cross validation includes leave-one-out, 10fold-cross-validation and 25fold-cross-validation and is performed during instantiation.

4.1.5. Dashboard

Just like the scanner, the dashboard module is only used by the recognizer app. It provides the user with a user interface showing charts and diagrams based on data by KPMG. That data is read and parsed by the dashboard when it is instantiated. As the user interface provides interactive elements like buttons the dashboard module is also responsible for handling user interaction. In addition, it exposes an interface to the recognition module which is used to broadcast recognized gestures as events to the dashboard. The dashboard can then execute the appropriate action.

4.1.6. Recognition

The recognition module used in the recognizer app is responsible to instantiate and hooking up the input module. It also instantiates the dashboard and implements the logic required to notify it about a recognized gesture.

5. Data Collection & Manipulation

Working with machine learning algorithms requires certain actions in order to choose, prepare and classify data streams. Choosing the right data is the key decision in order to be able to train the algorithm properly. In addition to this step, preparing data and transforming it into “features” is equally important.

5.1. Data Collection

To define which data will be collected we had to think about which kind of gestures we will want to recognize. Attaching the sensors to the back of our hand enables us to observe acceleration data for x, y, and z axis as well as tilt angles alpha, beta and gamma. We thought about attaching a magnet to our thumb so that we can use the hall-sensor as well, but we decided to leave this for further research. The collected acceleration and tilt angle data can later be used to distinguish gestures from each other. To avoid recording and scanning data observed during normal gesticulation we require the acceleration vector length of each gesture to exceed a certain threshold. In our code we set this threshold to 1.2 G.

5.2. Feature Design

Before we can train our algorithm, we have to transform the recorded data into a set of features. Choosing and designing those features affects the performance of the gesture recognition algorithm heavily.

5.2.1. Preprocessing

Standardization, scaling, normalization and binarization are common preprocessing techniques. Though, as the acceleration vector and tilt angle sizes are key characteristics to distinguish gestures scaling, normalization and binarization are not applicable for our use case.

5.2.2. Features

ToDo

Add images of feature selection We first thought about using the variance as a feature, but early tests already revealed that the acceleration data often has a variance near zero and thus gives us almost no information about the executed gesture.

Instead, maximum and minimum acceleration per axis differ widely across gestures and are worth adding to our feature set. To give our classifier a idea of which minimum/maximum was the first, we put them in the right order, so if we have a maximum first, this will be the first feature and the minimum will become the second feature. For instance, swiping left first shows a negative minimum acceleration, e.g. -1.7G, and then a positive maximum acceleration, e.g. 1.2G, on the x axis. We now pick those values out of our buffer and

figure out which one was earlier, so we can put them in our feature vector in the right order. In this example, we end up with the tuple $(-1.7, 1.2)$.

To include the tilt angles into our feature set, we had to come up with a different approach. Not only was their variance near zero, but also their minimum and maximum values are often equally great across different gestures. Summing up all angles led to non-satisfying results, so we collect now two features per angle: One sums up the absolute values of negative angle differences and the other one sums up the absolute values of the positive angle differences. Comparing both approaches the latter one increased our average precision score by 6%. **Include stats “First Approach” vs “Final Solution” here**

ToDo

5.3. Data Classification

Depending on the underlying training set, each classifier performs a bit different. We tested the `KNeighborsClassifier` with five and eleven nearest neighbors, `RandomForestClassifier` with 10 and 100 estimators, `DecisionTreeClassifier`, `SupportVectorMachine` and `GaussianNB`, which is a neural network. We evaluated each classifier with leave-one-out, 10fold-cross-validation and 25fold-cross-validation and summarized the results in the diagram below:

As you can see, the `GaussianNB` and `11NearestNeighbors` perform poorly compared to the other classifiers.

Appendix

A. First Appendix Section

...

